

JonJacobsonDSC630Exercise3.2

March 30, 2025

Load Packages and Data

```
[48]: import re
from operator import itemgetter

import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
from scipy.stats import gaussian_kde
from sklearn.model_selection import train_test_split
import sklearn.metrics as metrics
from sklearn.linear_model import LinearRegression

dodgers = pd.read_csv("dodgers-2022.csv")
row_count, column_count = dodgers.shape
print("There are {} rows and {} columns. Here is a sample:".format(
    row_count, column_count))
dodgers.head()
```

There are 81 rows and 12 columns. Here is a sample:

```
[48]:   month  day  attend  day_of_week  opponent  temp  skies  day_night  cap  shirt  \
0     APR    10      56000      Tuesday  Pirates    67  Clear       Day    NO    NO
1     APR    11      29729    Wednesday  Pirates    58  Cloudy      Night   NO    NO
2     APR    12      28328    Thursday  Pirates    57  Cloudy      Night   NO    NO
3     APR    13      31601     Friday   Padres    54  Cloudy      Night   NO    NO
4     APR    14      46549    Saturday  Padres    57  Cloudy      Night   NO    NO

   fireworks  bobblehead
0           NO          NO
1           NO          NO
2           NO          NO
3          YES          NO
4           NO          NO
```

```
[49]: attend = dodgers[['attend']]
min_attend = attend.min().iloc[0]
max_attend = attend.max().iloc[0]
```


```

print("Minimum Attendence: ",min_attend," Maximum Attendence: ",max_attend,"")
    ↵Min Percentage: {:.2f}%.format(min_attend/max_attend*100))
# Use bins that show bars for each possible value.
bins = np.array([- .25, 0.25, 0.75, 1.25, 1.75, 2.25, 2.75, 3.25, 3.75])
attend.plot.hist(bins=20,alpha=0.5,title="Histogram of Attendance")

```

Minimum Attendence: 24312 Maximum Attendence: 56000 Min Percentage: 43.41%

[49]: <Axes: title={'center': 'Histogram of Attendance'}, ylabel='Frequency'>



Analysis of Attendance Histogram The Dodgers stadium looks like it has a seating capacity of 56k. Quite a few games maxed out the stadium. The stadium was always more than 43.41% full. If you exclude the games that maxed out the stadium, the remainder is close to a standard

distribution with quite a bit of variability.

Check for Unique Values

```
[50]: for column in dodgers.columns:
    unique_values = dodgers[column].unique()
    count_unique = len(unique_values)
    if count_unique < 10:
        print("{}: {}".format(column,unique_values))
    else:
        print("{} has {} unique values".format(column,count_unique))
```

```
month: ['APR' 'MAY' 'JUN' 'JUL' 'AUG' 'SEP' 'OCT']
day has 31 unique values
attend has 80 unique values
day_of_week: ['Tuesday' 'Wednesday' 'Thursday' 'Friday' 'Saturday' 'Sunday'
'Monday']
opponent has 17 unique values
temp has 32 unique values
skies: ['Clear' 'Cloudy']
day_night: ['Day' 'Night']
cap: ['NO' 'YES']
shirt: ['NO' 'YES']
fireworks: ['NO' 'YES']
bobblehead: ['NO' 'YES']
```

Convert Categorical Columns into Dummy Columns

```
[52]: def safe_column(input):
    result = str(input).strip()
    result = result.lower()

    # convert special characters and whitespace to "_"
    result = re.sub(r"[^a-zA-Z0-9_]", "_", result)

    # reduce adjacent "_" characters
    result = re.sub(r"_+", "_", result)

    result = re.sub(r"^\_", "", result) # remove leading "_"
    result = re.sub(r"\_$", "", result) # remove trailing "_"
    if not result:
        result = "none"
    return result

# Turn categorical columns into dummy columns.
def make_dummies(df,column):
    dummies = pd.get_dummies(df[column]).rename( \
        columns=lambda x: column + '_' + safe_column(x))
    df = pd.concat([df, dummies], axis=1)
```

```

df.drop([column], inplace=True, axis=1)
return df

with_dummies = dodgers.copy()
categorical = [
'month',
'day_of_week',
'opponent',
]
print("Shape before creating dummy columns: ", with_dummies.shape)
for column in categorical:
    with_dummies = make_dummies(with_dummies, column)
    print("Shape after converting "+column+" into dummy columns:", \
        with_dummies.shape)

```

Shape before creating dummy columns: (81, 12)
 Shape after converting month into dummy columns: (81, 18)
 Shape after converting day_of_week into dummy columns: (81, 24)
 Shape after converting opponent into dummy columns: (81, 40)

Convert Boolean Columns into Integers

[53]:

```

# Handle boolean columns.
with_dummies = with_dummies.rename(columns={
    'day_night': 'day_game',
    'skies':'clear_skies'
})
with_dummies['day_game'] = with_dummies['day_game'].map({
    'Day':'YES',
    'Night':'NO',
})
with_dummies['clear_skies'] = with_dummies['clear_skies'].map({
    'Clear ':'YES',
    'Cloudy':'NO',
})
with_dummies = with_dummies.replace({
    'YES':1,
    'NO':0
})

with_dummies.head()

```

[53]:

	day	attend	temp	clear_skies	day_game	cap	shirt	fireworks	\
0	10	56000	67	1	1	0	0	0	
1	11	29729	58	0	0	0	0	0	
2	12	28328	57	0	0	0	0	0	
3	13	31601	54	0	0	0	0	1	
4	14	46549	57	0	0	0	0	0	

```

    bobblehead month_apr ... opponent_marlins opponent_mets \
0          0     True ...             False      False
1          0     True ...             False      False
2          0     True ...             False      False
3          0     True ...             False      False
4          0     True ...             False      False

    opponent_nationals opponent_padres opponent_phillies opponent_pirates \
0            False        False        False      True
1            False        False        False      True
2            False        False        False      True
3            False        True         False      False
4            False        True         False      False

    opponent_reds opponent_rockies opponent_snakes opponent_white_so
0            False        False        False      False
1            False        False        False      False
2            False        False        False      False
3            False        False        False      False
4            False        False        False      False

[5 rows x 40 columns]

```

Check for Correlations

```
[54]: # Adapted from https://www.geeksforgeeks.org/
       ↵create-a-correlation-matrix-using-python/
import matplotlib.pyplot as plt
import matplotlib.colors as colors
import numpy as np
from sklearn import datasets
import pandas as pd

matrix = with_dummies.corr()

# Mask the duplicated upper triangle and the diagonal
mask = np.triu(np.ones_like(matrix, dtype=bool))
masked_matrix = np.ma.masked_array(matrix, mask)

# plot the masked correlation matrix
plt.imshow(masked_matrix, vmin=-1, vmax=1, cmap="Blues")

# adding colorbar
plt.colorbar()

# extracting variable names
variables = []

```

```


for i in matrix.columns:
    variables.append(i)

# Adding labels to the matrix
plt.xticks(range(len(matrix)), variables, rotation=45, ha='right')
plt.yticks(range(len(matrix)), variables)

# Make the figure full screen
plt.rcParams['figure.figsize'] = [24, 24]

# Display the plot
plt.show()

```



```
[55]: # Set a threshold for strong correlations
threshold = 0.5

# Iterate through the correlation matrix and filter based on the threshold
correlations = []
seen = {}

for column in matrix.columns:
    for index in matrix.index:
        seen_key = [column, index]
        seen_key.sort()
        seen_key = ''.join(seen_key)
        if seen_key in seen:
            continue
        else:
            seen[seen_key] = 1
        if index != column and abs(matrix.loc[index, column]) > threshold:
            correlation = matrix.loc[index, column]
            correlations.append({
                'correlation': correlation,
                'index': index,
                'column': column,
            })

cors = sorted(correlations, key=lambda d: d['correlation'])
for cor in cors:
    print("{:+.5f} {} -> {}".format(cor['correlation'], cor['index'], cor['column']))
```

```
-0.51200 month_apr -> temp
+0.55470 opponent_giants -> month_oct
+0.55470 opponent_angels -> month_jun
+0.55470 opponent_white_sox -> month_jun
+0.56011 month_sep -> temp
+0.58189 bobblehead -> attend
+0.74399 day_of_week_sunday -> day_game
+0.95651 day_of_week_friday -> fireworks
```

Analyze Correlations There are some correlations here that are expected, such as the temperature correlating to the month. Other correlations show patterns in programming:

- Fireworks are generally done on Fridays
- Sundays are generally day games
- The dodgers played the giants in October and played the Angels and White Sox in Jun.

There is a critical correlation here that might provide an indication of how to improve attendance, which is the “bobblehead” condition. When the “bobblehead” is present, attendance tends to be higher. Whether this proves causation or is similar to programming decisions, such as having fireworks on Fridays, would require further analysis.

Attempt a Linear Regression Model

```
[56]: # Split the data into predictor data and target column
target_column = 'attend'
data = with_dummies.copy()
target = data[target_column]
data.drop([target_column], inplace=True, axis=1)

# Split the data into 80% training and 20% testing
data_train, data_test, target_train, target_test = train_test_split(data, target, test_size=0.2, random_state=42)

print("Training Data: ", data_train.shape, "Test data: ", data_test.shape)
```

Training Data: (64, 39) Test data: (17, 39)

```
[57]: def print_model_metrics(test,pred):
    mae = metrics.mean_absolute_error(test, pred)
    mse = metrics.mean_squared_error(test, pred)
    rmse = np.sqrt(mse) # or mse**(.5)
    r2 = metrics.r2_score(test, pred)

    print("R-Squared:", r2)
    print("Root Mean Squared Error (RMSE):", rmse)
    print("Mean Absolute Error (MAE): ",mae)

# Build the model with the columns in a specific order, so that
# we can pair the columns with the coefficients for analysis later.
model_columns = data_train.columns
model = LinearRegression().fit(data_train[model_columns],target_train)
coefficients = list(zip(model_columns,model.coef_))

# Obtain the predictions
target_pred = model.predict(data_test)

print("LinearRegression Metrics:")
print_model_metrics(target_test,target_pred)
```

LinearRegression Metrics:

R-Squared: 0.12209095246024049

Root Mean Squared Error (RMSE): 9290.744150802151

Mean Absolute Error (MAE): 7386.547606778323

Coefficients

```
[58]: sorted_coef = sorted(coefficients, key=itemgetter(1))
for coef in sorted_coef:
    print(coef[0],coef[1])
```

day_of_week_friday -16118.29583293249

month_apr -11585.856321502955

```
cap -7589.533374940717
opponent_astros -6480.1799221782685
opponent_marshallins -6315.423496325919
opponent_reds -3317.3756456068495
opponent_snakes -3130.913350932372
opponent_giants -3105.239462902717
month_may -2577.406525736656
day_of_week_monday -2437.162946142595
shirt -2353.8821784658508
opponent_rockies -1945.2329553281745
opponent_brewers -1928.98173640902
clear_skies -766.3445584847514
month_sep -375.4398870130925
opponent_white_sox -313.1509038970658
day_game -269.3709733047583
temp -181.81774654577
day_of_week_wednesday -167.0017138577736
day 69.02131415946556
opponent_pirates 160.24885651932175
day_of_week_thursday 307.90277894228194
opponent_mets 769.0588089288347
opponent_phillies 1128.1672731795445
opponent_cubs 1373.9951151149714
opponent_cardinals 1499.2565492629715
opponent_angels 1611.040269508187
month_jul 1688.4845466537483
opponent_braves 1848.4947800042069
opponent_padres 1921.6672560444722
day_of_week_saturday 3447.943263614855
month_oct 3734.6884648784453
month_jun 4056.3856913385575
month_aug 5059.144031381842
day_of_week_tuesday 7448.336531106133
day_of_week_sunday 7518.277919269568
bobblehead 9672.400714416563
opponent_nationals 16224.568565017835
fireworks 18538.962126872368
```

Analysis of Linear Regression There is so little data, that the use of a linear regression model is risky. And, as it turns out, the Linear Regression was not particularly successful. An R-Squared value of 0.12 means that very little of the variance in attendance can be explained by the linear model. Of the variance that was explained, Tuesday and Sunday games were more likely to increase attendance. Also, fireworks and bobblehead condition increased attendance. Lastly, having the Nationals as the opponent also increased attendance. However, based on the low r-squared value and the limited data, these insights should be used cautiously.

Conclusion This analysis would benefit from significantly more data, encompassing more years and more teams. With what is available, however, the best single recommendation I could make to increase attendance would be to increase the occurrence of the bobblehead condition. I need more domain knowledge on this characteristic to know how best to make use of this insight, or even if it applies. It could be that some other factor is independently driving both the use of the bobblehead and attendance. However, this is the variable that correlates most closely with attendance, so it is a good place to focus. It might also be worth considering playing more games against the Nationals to take advantage of the increased attendance that opponent appears to give. Lastly, it may be worth it to consider using fireworks more, but this would need to be weighed against the costs of adding additional fireworks displays.

[]: