

# JacobsonDSC550Week4

December 22, 2024

In the Week 4 Exercise, you will build a linear regression model to predict fuel efficiency (miles per gallon) of automobiles. Download the auto-mpg.csv dataset from: Auto-mpg dataset. Load the data as a Pandas data frame and ensure that it imported correctly.

```
[30]: import pandas as pd
mpg = pd.read_csv("auto-mpg.csv")
rows,columns=mpg.shape
print("There are {} rows and {} columns.".format(rows,columns))
mpg[mpg['origin']==3].head()
```

There are 398 rows and 9 columns.

```
[30]:    mpg  cylinders  displacement horsepower  weight  acceleration \
14   24.0          4        113.0         95    2372       15.0
18   27.0          4         97.0         88    2130       14.5
29   27.0          4         97.0         88    2130       14.5
31   25.0          4        113.0         95    2228       14.0
53   31.0          4         71.0         65    1773       19.0

      model  year  origin           car name
14            70     3  toyota corona mark ii
18            70     3      datsun pl510
29            71     3      datsun pl510
31            71     3  toyota corona
53            71     3  toyota corolla 1200
```

Begin by prepping the data for modeling:

- \* Remove the car name column.
- \* The horsepower column values likely imported as a string data type. Figure out why and replace any strings with the column mean.
- \* Create dummy variables for the origin column. (It appears that 1 Is American, 2 is European, 3 is Japanese)

```
[31]: mpg = mpg.drop(columns=['car name'])

# Replace string horsepower values with the mean value for the column
horsepower_mean = pd.to_numeric(mpg['horsepower'], errors='coerce').dropna().mean()
mask = pd.to_numeric(mpg['horsepower'], errors='coerce').isna()
```

```

print("These are the non-numeric values in the 'horsepower' column, which
    ↪caused it to be read as a string:\n",mpg['horsepower'][mask],"\\nWe are
    ↪replacing them with the mean of: ",horsepower_mean)
modified = mpg['horsepower'].copy()
modified[mask]=horsepower_mean
mpg['horsepower']=modified
mpg['horsepower'] = pd.to_numeric(mpg['horsepower'])

# Add dummy variables (binaries) for origin values
mpg['american']=mpg['origin']==1
mpg['european']=mpg['origin']==2
mpg['japanese']=mpg['origin']==3

```

These are the non-numeric values in the 'horsepower' column, which caused it to be read as a string:

```

32      ?
126     ?
330     ?
336     ?
354     ?
374     ?

Name: horsepower, dtype: object
We are replacing them with the mean of: 104.46938775510205

```

Create a correlation coefficient matrix and/or visualization. Are there features highly correlated with mpg?

[32]: mpg.corr()

	mpg	cylinders	displacement	horsepower	weight	\
mpg	1.000000	-0.775396	-0.804203	-0.771437	-0.831741	
cylinders	-0.775396	1.000000	0.950721	0.838939	0.896017	
displacement	-0.804203	0.950721	1.000000	0.893646	0.932824	
horsepower	-0.771437	0.838939	0.893646	1.000000	0.860574	
weight	-0.831741	0.896017	0.932824	0.860574	1.000000	
acceleration	0.420289	-0.505419	-0.543684	-0.684259	-0.417457	
model year	0.579267	-0.348746	-0.370164	-0.411651	-0.306564	
origin	0.563450	-0.562543	-0.609409	-0.453669	-0.581024	
american	-0.568192	0.604351	0.651407	0.486083	0.598398	
european	0.259022	-0.352861	-0.373886	-0.281258	-0.298843	
japanese	0.442174	-0.396479	-0.433505	-0.321325	-0.440817	
	acceleration	model year	origin	american	european	japanese
mpg	0.420289	0.579267	0.563450	-0.568192	0.259022	0.442174
cylinders	-0.505419	-0.348746	-0.562543	0.604351	-0.352861	-0.396479
displacement	-0.543684	-0.370164	-0.609409	0.651407	-0.373886	-0.433505
horsepower	-0.684259	-0.411651	-0.453669	0.486083	-0.281258	-0.321325
weight	-0.417457	-0.306564	-0.581024	0.598398	-0.298843	-0.440817

acceleration	1.000000	0.288137	0.205873	-0.250806	0.204473	0.109144
model year	0.288137	1.000000	0.180662	-0.139883	-0.024489	0.193101
origin	0.205873	0.180662	1.000000	-0.924486	0.246332	0.886596
american	-0.250806	-0.139883	-0.924486	1.000000	-0.597198	-0.643317
european	0.204473	-0.024489	0.246332	-0.597198	1.000000	-0.229895
japanese	0.109144	0.193101	0.886596	-0.643317	-0.229895	1.000000

```
[34]: # Adapted from https://www.geeksforgeeks.org/
       ↪create-a-correlation-matrix-using-python/
import matplotlib.pyplot as plt
from sklearn import datasets
import pandas as pd
matrix = mpg.corr()

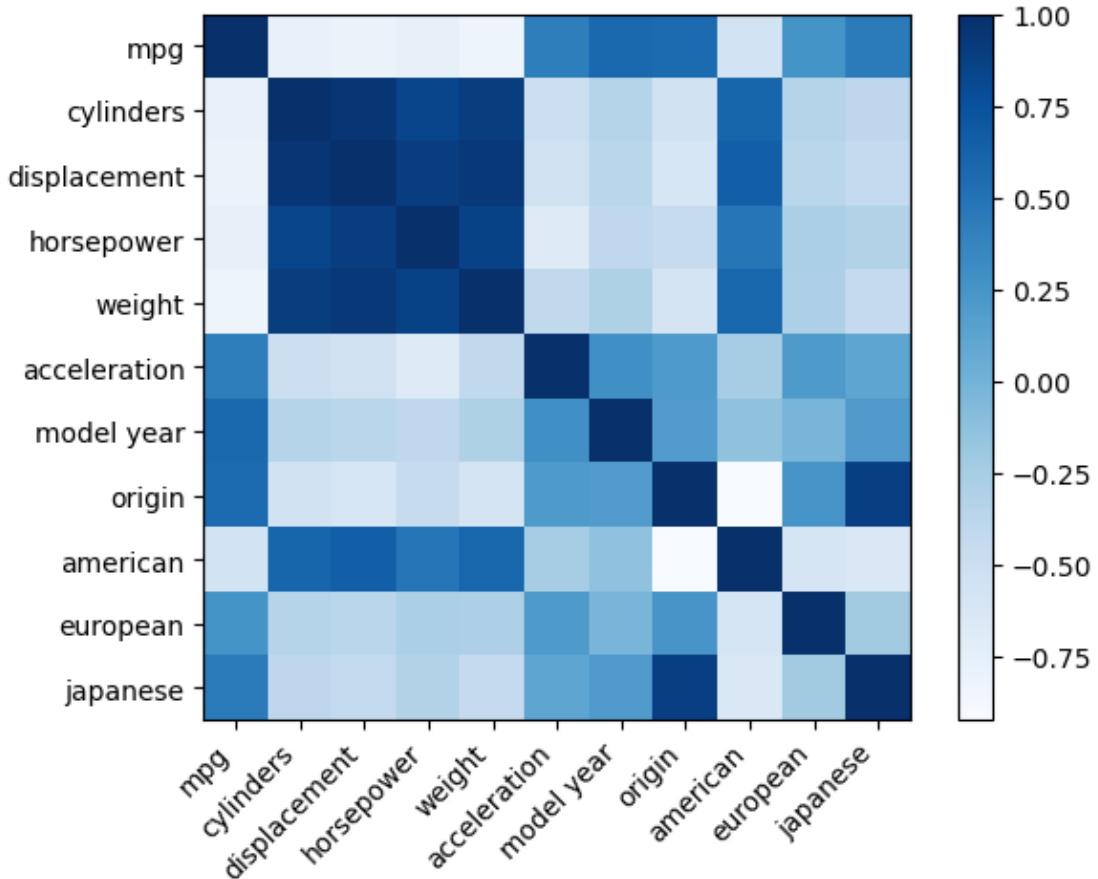
# plotting correlation matrix
plt.imshow(matrix, cmap='Blues')

# adding colorbar
plt.colorbar()

# extracting variable names
variables = []
for i in matrix.columns:
    variables.append(i)

# Adding labels to the matrix
plt.xticks(range(len(matrix)), variables, rotation=45, ha='right')
plt.yticks(range(len(matrix)), variables)

# Display the plot
plt.show()
```



Looking for correlation values greater than 0.5, the following values are all closely correlated with each other: \* cylinders \* displacement \* horsepower \* weight

These same values are inversely correlated with mpg.

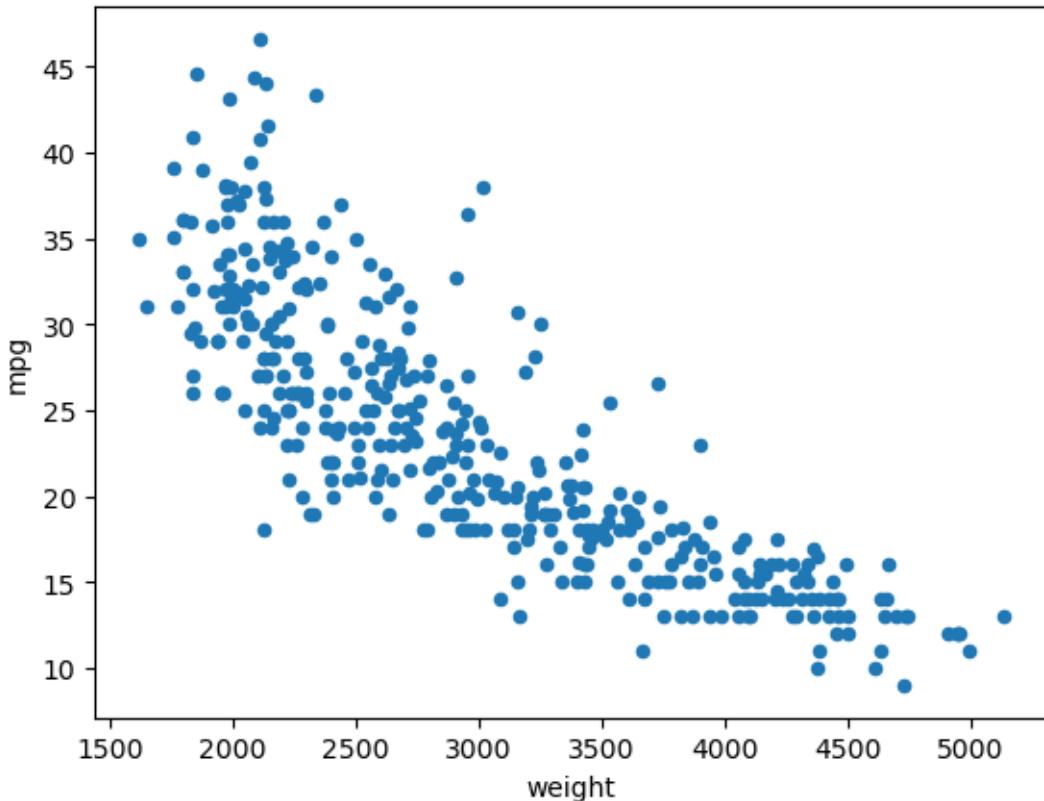
Also, american is correlated with cylinders and displacement, and inversely correlated with mpg.

Additionally, mpg and model year are correlated.

**Plot mpg versus weight. Analyze this graph and explain how it relates to the corresponding correlation coefficient.**

```
[36]: mpg.plot.scatter(x='weight', y='mpg')
```

```
[36]: <Axes: xlabel='weight', ylabel='mpg'>
```



Weight and mpg are inversely correlated fairly strongly with a coefficient of -0.831741. You can see this on the graph in the fact that a trend line would slope downwards towards the right. You can also see that the pattern is a downward slope and might be modeled better logarithmically, but is still modeled very well linearly. Mpg clearly decreases, on average, as weight increases.

**Randomly split the data into 80% training data and 20% test data, where your target is mpg.**

```
[38]: from sklearn.model_selection import train_test_split

data = mpg.drop('mpg', axis=1)
target = mpg['mpg']

# Split the data into 80% training and 20% testing
data_train, data_test, target_train, target_test = train_test_split(data, target, test_size=0.2, random_state=42)

print(data_train.shape, data_test.shape)
print(target_train.shape, target_test.shape)
```

(318, 10) (80, 10)  
(318,) (80,)

Train an ordinary linear regression on the training data.

```
[42]: from sklearn.linear_model import LinearRegression
linear_model = LinearRegression()
linear_model.fit(data_train,target_train)
print(linear_model.intercept_)
print(linear_model.coef_)

# Obtain the predictions
target_pred = linear_model.predict(data_test)
```

-21.971261257955092  
[-0.16373048 0.01958399 -0.01334457 -0.00707275 0.07335016 0.82739747  
0.88430052 -0.97974802 1.07519552 -0.0954475 ]

Calculate R<sup>2</sup>, RMSE, and MAE on both the training and test sets and interpret your results.

```
[55]: import sklearn.metrics as metrics
import numpy as np

def print_model_metrics(test,pred):
    mae = metrics.mean_absolute_error(test, pred)
    mse = metrics.mean_squared_error(test, pred)
    rmse = np.sqrt(mse) # or mse**(0.5)
    r2 = metrics.r2_score(test, pred)

    print("R-Squared:", r2)
    print("Root Mean Squared Error (RMSE):", rmse)
    print("Mean Absolute Error (MAE): ",mae)

print("Linear Regression Model Metrics:")
print_model_metrics(target_test,target_pred)
```

Linear Regression Model Metrics:  
R-Squared: 0.8449006123776615  
Root Mean Squared Error (RMSE): 2.8877573478836327  
Mean Absolute Error (MAE): 2.2875867704421085

Pick another regression model and repeat the previous two steps. Note: Do NOT choose logistic regression as it is more like a classification model.

```
[56]: from sklearn.preprocessing import PolynomialFeatures

# Train polynomial regression model on the whole dataset
poly = PolynomialFeatures(degree = 2, include_bias=False)
train_poly = poly.fit_transform(data_train)
test_poly = poly.fit_transform(data_test)
```

```
poly_model = LinearRegression()
poly_model.fit(train_poly, target_train)
poly_pred=poly_model.predict(test_poly)

print("Polynomial Regression Model Metrics:")
print_model_metrics(target_test,poly_pred)
```

Polynomial Regression Model Metrics:  
R-Squared: 0.8945181878762949  
Root Mean Squared Error (RMSE): 2.3814663378257643  
Mean Absolute Error (MAE): 1.7498402349944164

[ ]: